

Identifying myocardial infarction risk factors in the Wisconsin Longitudinal Survey to aid in intervention program design

Fred Boehm, Statistics 998

March 31, 2015

Todo list

■ What is mechanism for smoking causing CAD?	1
■ what are other known risk factors?	1
■ need more info on WLS?	2
■ what to explore for categorical data? Only the proportion in each category????	2
■ Show a single ROC plot & explain what is meant by "AUC"	4
■ explain why we want Loh's algorithms. Might mention why they could be useful in this project.	4

Abstract

Introduction

Coronary artery disease (CAD) is a leading cause of death in the United States and much of North America and Europe. In 2011, one American died of CAD every 40 seconds, on average, and 155,000 of those deaths were people aged less than 65 years.¹ One manifestation of CAD is a myocardial infarction (MI), which is also called a “heart attack”. A MI results from a clot in a coronary artery that diminishes blood flow to the heart muscle, or myocardium. If blood flow disruption persists for a sufficiently long time, the muscle may die, or infarct. The irreparable dead heart muscle diminishes the overall ability of the heart to pump blood. Severe MIs may lead to a patient’s death.

Epidemiologists have identified modifiable and non-modifiable risk factors that contribute to CAD risk. Smoking is among the strongest modifiable risk factors, and is thought to elevate CAD risk by triggering elevations in inflammatory molecules in the bloodstream. _____

What is mechanism for smoking causing CAD?

. Diabetes mellitus and hypertension (systolic or diastolic) are typically considered non-modifiable risk factors, although their contribution to CAD risk may be reduced in patients who undertake dramatic lifestyle interventions, such as exercise programs and diet with weight loss. Non-modifiable risk factors include age, a family history of CAD and presence of certain genetic variants

what are other known risk factors?

.
Our collaborators at the Wisconsin Longitudinal Study (WLS) have undertaken an investigation on a subset of WLS participants with the goal of identifying CAD risk factors in the WLS study population.

¹Mozaffarian et al., “Executive Summary.”

The ultimate goal of this project is to develop an intervention program to reduce CAD morbidity and mortality in Wisconsin. The investigators would like to extend such an intervention program to Wisconsin residents who are not WLS subjects. Our goal in this report is to identify risk factors for MI among WLS participants.

Study design

The Wisconsin Longitudinal Study (WLS) is a long-term study of a random sample of 10,317 men and women who graduated from Wisconsin high schools in 1957. According to the WLS website “WLS provides an opportunity to study the life course, intergenerational transfers and relationships, family functioning, physical and mental health and well-being, and morbidity and mortality from late adolescence through 2011.”²

need more info on WLS?

Our collaborators collected data from the original respondents or their parents in 1957, 1964, 1975, 1992, 2004, and 2011; from a selected sibling in 1977, 1994, 2005, and 2011; from the spouse of the original respondent in 2004; from the spouse of the selected sibling in 2006; and from widow(er)s of the graduates and siblings in 2006.

Data description

Our collaborators shared with us a data set that contains records for 19095 individuals (including original subjects and siblings) with 310 variables per subject. 9363 subjects responded (with yes or no) to the 2011 question of whether they had ever had a heart attack.

Exploratory data analyses

Since our WLS data included 310 variables, we won't provide summaries for all of them in this document. Instead, we focus our exploratory analyses on our response variables (HAC2011, HAC2004, HACinc, doc2011, doc2004, docinc) and covariates that other researchers have identified as associated with coronary artery disease.

Age is known, from epidemiologic studies, to be a strong risk factor for CAD, with older individuals having an elevated CAD risk.

what to explore for categorical data? Only the proportion in each category????

Statistical modeling

We used statistical modeling to try to identify covariates that associated with six distinct outcomes: 1) HAC2004, 2) HAC2011, 3) DOC2004, 4) DOC2011, 5) new self-reported heart attacks (from 2004 to 2011) and 6) new heart attack per doctor's report (from 2004 to 2011). For a subject to qualify as a “new” self-reported heart attack, they must have responded “No” in 2004 and “Yes” in 2011. Analogous definition applies for “new” doctor-reported heart attack.

We found that HAC2004 had 11534 non-missing responders (with 665 responding “Yes” and 10869 responding “No”). Counts for other variables are provided in Appendix A.

²“Wisconsin Longitudinal Study”.

Framingham study variable	WLS Variable
Sex	Sex
Quantitative total cholesterol	highchol2011, highchol2004
Quantitative HDL cholesterol	None
Smoking	smokever2011 (Columns 61 to 87 aim to quantify smoking)
Diabetes	diabetes2004, diabetes2011
Age	Age
Systolic BP	highbp2004, highbp2011*
Treated for high blood pressure	None

Table 1: Framingham study variables and their closest analogs in WLS. (* SBP not available, so we used reported "high BP".)

Statistical modeling with Framingham study predictors

We identified those variables in the WLS that closely match those in the Framingham study³ (Table ??). It's important to note that the Framingham study used survival analysis methods, including Cox proportional hazards regression, to identify risk factors for a cardiovascular event. Thus, their study design, analysis, and purpose differ from ours.

Because the risk factors from the Framingham study have strong effect sizes, are easily interpreted, and widely used by both physicians and public health scientists, we decided to perform logistic regression analyses with solely those WLS variables that most closely matched the Framingham variables. We analyzed each of the 6 outcomes of interest. For each outcome variable, we fitted a logistic regression model using the entire data set (omitting those subjects with missing data). We then transformed the fitted logit values to probabilities before plotting a receiver operating characteristic (ROC) curve for each model, in which we use the fitted probabilities to examine the trade-off between specificity and sensitivity. We also calculated the area under the curve (AUC) for each ROC curve.⁴

Subgroup-specific modeling

Results

Show a single ROC plot & explain what is meant by "AUC"

Comparing statistical models with cross-validation and area under the curve (AUC) statistic

In evaluating our statistical models, we used “area under the curve” (AUC) from our receiver operating characteristic (ROC) curves. Each statistical model was evaluated with 5-fold cross-validation. For those unfamiliar with cross-validation, five-fold cross-validation means that we partition the subjects into 5 non-overlapping “folds” of approximately equal size. We then fit 5 models, using 4 of the 5 folds as the “training” set for each fit. For each fit, we then test the model with fold that wasn’t used in the corresponding training set. For example, for five-fold CV, we first omit fold #1 and fit the model using folds 2,3,4, and 5 together. We then test the model using fold #1. We then fit a model with folds 1,3,4, and 5 together and test it with fold #2. We continue this procedure until all five possible models are fitted.

³D’Agostino et al., “General Cardiovascular Risk Profile for Use in Primary Care the Framingham Heart Study.”

⁴Robin et al., “PROC.”




Figure 1: Receiver operating characteristic curve for the logistic regression model with outcome HAC2011 and the six Framingham covariates.

For each of the five models, we created a received operating characteristic (ROC) curve and evaluated the area under the curve using the pROC R package.⁵

A ROC curve is a plot of sensitivity against specificity. We created the ROC curve by varying the threshold cut point for our classifier and seeing how distinct cut points impact our specificity and sensitivity as measured on the test set.

Discussion

Future directions

We would like to investigate other tree-based methods, including those developed by Wei-Yin Loh's research team. The absence of R implementations of Loh's algorithms prevented us from including them in our current analysis. We have begun the process of writing code to implement Loh's algorithms in R, but they are not yet ready for use.

explain why we want Loh's algorithms. Might mention why they could be useful in this project.

A major reason for using Loh's algorithms, such as GUIDE, is their unbiasedness.⁶ CART, the algorithm that's implemented in the rpart package, has an intrinsic selection bias that favors selection of categorical variables with more discrete values. GUIDE, on the other hand, avoids this selection bias. In the current scenario, because most of our variables are either continuous or binary, we're not selection bias is likely to be a major problem, but we'd still like to compare GUIDE results with those of CART.

⁵Ibid.

⁶Loh, "Classification and Regression Trees."

Given the relatively low cost of acquiring genomics data, one possible future direction is to acquire genomics data for a subset of study subjects. For example, SNP genotype data from each subject may enable us to further discriminate those subjects that are at high risk for a CAD event. In some cases, we may be able to use cheek swabs as sources of DNA, which would enable data collection by mail. Such knowledge of genetic risk factors, when coupled with non-genetic risk factors, may be translated into the proposed intervention program, for example, by promoting healthy diet and physical activity among those at greatest risk.

Appendix A: Supplementary materials

Appendix B: Questions for client

1. What do you intend to do with results from this study?
2. How important is prediction of future MI to your scientific goals? There may be a tradeoff between prediction and the ability to quantify variable effects (for example does smoking increase or decrease risk? By how much? Does the effect vary across population subgroups?).
3. Why are there so many missing values (NA) for the HA2011 & HA2004 variables? Did these subjects not respond to that question? Did they respond to the survey at all?
4. Which variables do you think are the most meaningful?
5. Is incidence of MI between 2004 and 2011 a meaningful outcome? We could try to ascertain who had a MI during that interval (from among the people who had no MI history in 2004)
6. We're considering focusing on established risk factors such as those that the Framingham study identified. Is this reasonable? Or do you want to try to identify novel risk factors?
7. Are the "created" variables HAC2011 and HAC2004 better than the unadjusted versions? Is there any reason to do separate analyses for 2011 and 2004? Could we just code "yes" for a yes in any year, otherwise use 2011 value.
Why do some responses switch from "yes" in 2004 to "no" in 2011—is this because the question specifically refers to heart attacks or diagnoses within the last 20 years?
8. Should any conditions be excluded as covariates because they can occur concurrently with heart disease (e.g. diabetes, stroke, high blood pressure, high cholesterol)? In other words, is there more interest in leading indicators?
9. Is there interest in separating the effects of covariates gathered in multiple years (e.g. highchol2004, highchol2011)? Should these be combined into a single measure? Should 2011 effect be excluded when modeling a 2004 response?
10. I noticed some ordinal variables (e.g. education level) are coded as integers. Is this consistent for all ordinal variables? Are there any non-ordinal categorical variables coded as integers?

Appendix C: Computing code

Appendix D: Additional note

Throughout this report, we tried to adhere to the style suggested by Leek.⁷ We used the R statistical environment for all calculations⁸. Mike Wurm performed the tree-based analysis, and the code for that part is nearly identical to his. The R packages gbm,⁹ pROC¹⁰ and rpart¹¹ played central roles in our analyses.

References

- D’Agostino, Ralph B, Ramachandran S Vasan, Michael J Pencina, Philip A Wolf, Mark Cobain, Joseph M Massaro, and William B Kannel. “General Cardiovascular Risk Profile for Use in Primary Care the Framingham Heart Study.” *Circulation* 117, no. 6 (2008): 743–53.
- Leek, Jeff. *The Elements of Data Analytic Style*. Leanpub, 2015 Available at <https://leanpub.com/datastyle>.
- Loh, Wei-Yin. “Classification and Regression Trees.” *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery* 1, no. 1 (2011): 14–23.
- Mozaffarian, Dariush, Emelia J Benjamin, Alan S Go, Donna K Arnett, Michael J Blaha, Mary Cushman, Sarah de Ferranti, et al. “Executive Summary: Heart Disease and Stroke Statistics—2015 Update a Report from the American Heart Association.” *Circulation* 131, no. 4 (2015): 434–41.
- others, Greg Ridgeway with contributions from. *Gbm: Generalized Boosted Regression Models*, 2015. <http://CRAN.R-project.org/package=gbm> R package version 2.1.1.
- R Core Team. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing, 2015. <http://www.R-project.org/>.
- Robin, Xavier, Natacha Turck, Alexandre Hainard, Natalia Tiberti, Frédérique Lisacek, Jean-Charles Sanchez, and Markus Müller. “PROC: An Open-Source Package for R and S+ to Analyze and Compare ROC Curves.” *BMC Bioinformatics* 12 (2011): 77.
- Therneau, Terry, Beth Atkinson, and Brian Ripley. *Rpart: Recursive Partitioning and Regression Trees*, 2015. <http://CRAN.R-project.org/package=rpart> R package version 4.1-9.
- “Wisconsin Longitudinal Study”. <http://www.ssc.wisc.edu/wlsresearch/>, 2015 Accessed on 21-March-2015.

⁷*The Elements of Data Analytic Style*.

⁸R Core Team, *R*.

⁹Others, *Gbm*.

¹⁰Robin et al., “PROC.”

¹¹Therneau, Atkinson, and Ripley, *Rpart*.